

# **ICS 2305 : SYSTEMS PROGRAMMING**

---

NETWORKING

Karanja Mwangi

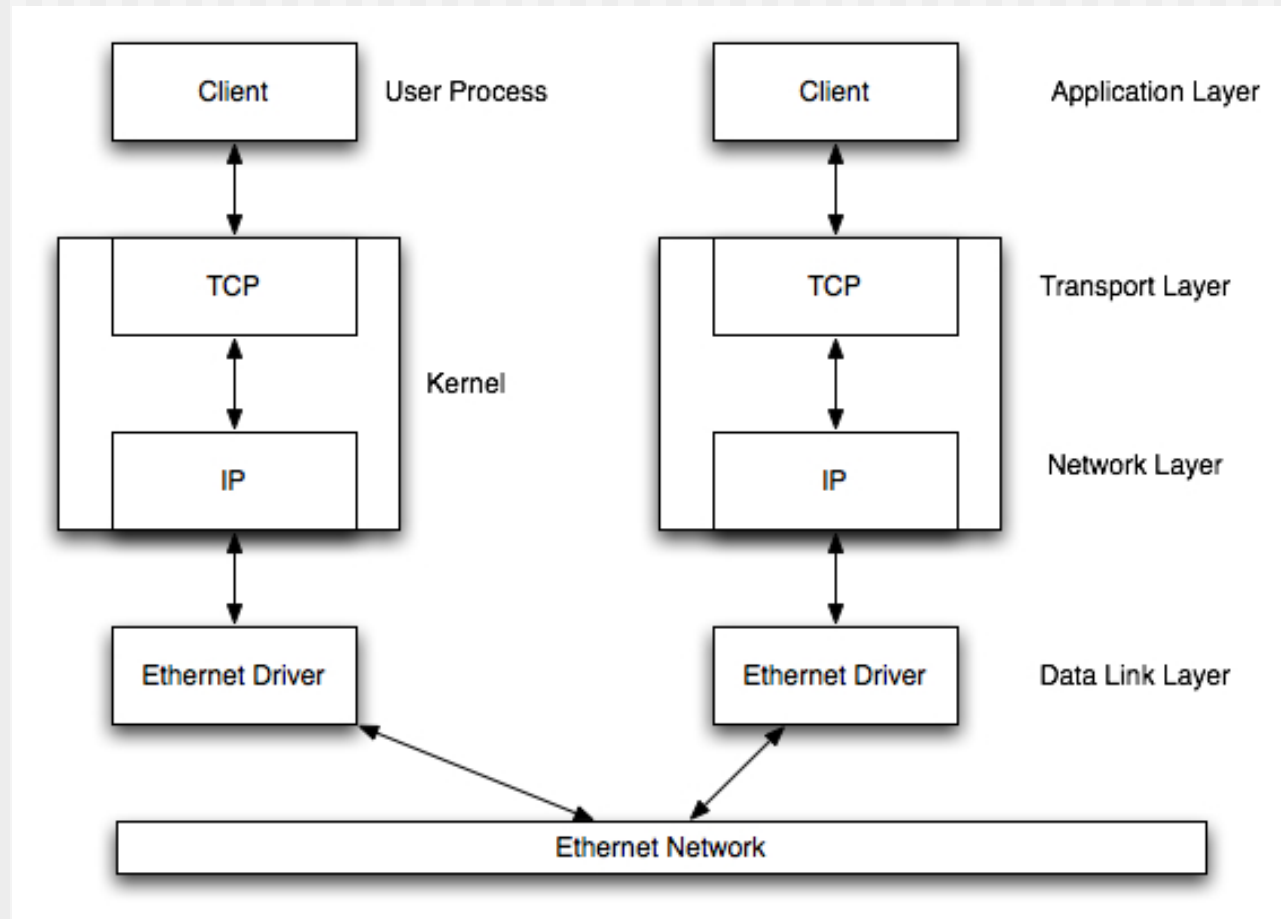
## Lesson objective

---

- **At the end of this class you will**
  - Refresh on Networking concepts and how they relate to systems programming
  - Able to build your client and server (TCP) applications in C
  - Use C programs to interact with Ports

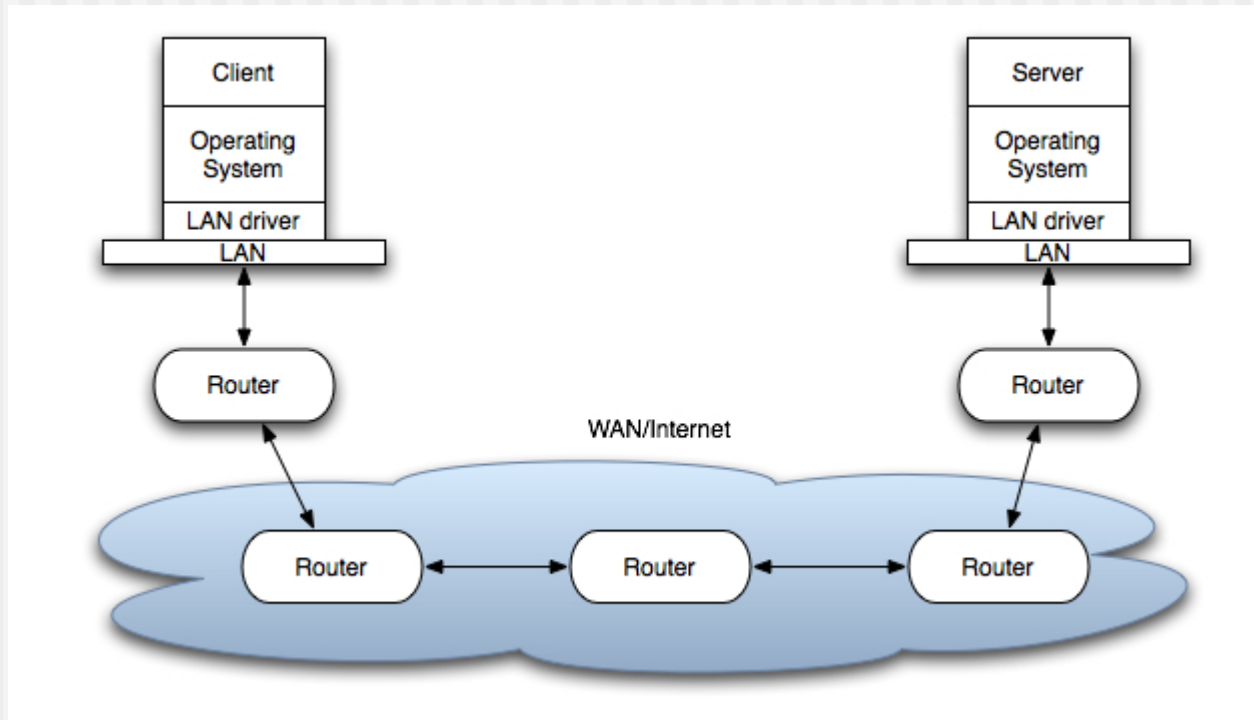
## Same LAN

client and the server on the same local network



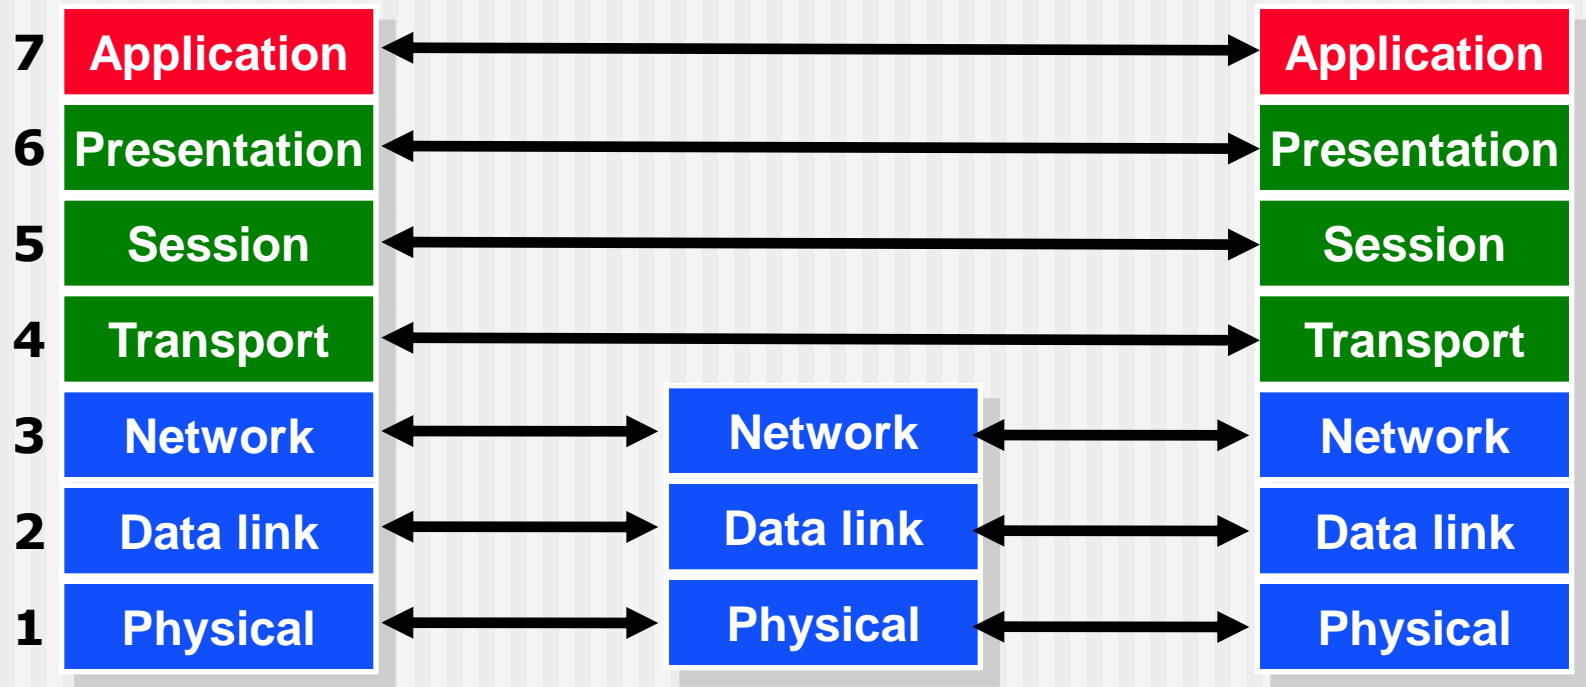
# TCP/IP

The client and the server may be in different LANs, with both LANs connected to a Wide Area Network (WAN) by means of *routers*.



# Network Layering

---



# Client Server Communication TCP and UDP

- The main transport protocols TCP and UDP enables communication between network applications
- UDP is a **connectionless protocol** (**RFC 768**)
  - built on **top of IPv4 and IPv6**.
  - Determines the **destination address** and **port** and send your data packet!
  - Packets or Datagrams may be dropped if the network is congested. Packets may be duplicated or arrive out of order.--- **No guarantee on good delivery**

# Client Server Communication-TCP

TCP is a **connection-based protocol** that is built on top of IPv4 and IPv6 ("TCP/IP" or "TCP over IP"). source IP address and source port number.

- TCP -**Manages the packets**, re-arranging out-of-order packets and changing the rate at which packets are sent, removes the duplicate etc
- TCP manages packets through 3 handshake **SYN, SYN-ACK, and ACK**
- For more on handshaking read [https://www.inetdaemon.com/tutorials/internet/tcp/3-way\\_handshake.shtml](https://www.inetdaemon.com/tutorials/internet/tcp/3-way_handshake.shtml)
  - RFC 793, RFC 1323, RFC 2581 and RFC 3390.

# User Datagram Protocol(UDP): An Analogy

## UDP

- **Single socket to receive messages**
- **No guarantee of delivery**
- **Not necessarily in-order delivery**
- **Datagram – independent packets**
- **Must address each packet**

## Postal Mail

- **Single mailbox to receive letters**
- **Not very reliable**
- **Not necessarily in-order delivery**
- **Letters sent independently**
- **Must address each reply**

Example UDP applications

**Multimedia, voice over IP**



# Transmission Control Protocol (TCP): An Analogy

## TCP

- **Reliable – guarantee delivery**
- **Byte stream – in-order delivery**
- **Connection-oriented – single socket per connection**
- **Setup connection followed by data transfer**

## Telephone Call

- **Guaranteed delivery**
- **In-order delivery**
- **Connection-oriented**
- **Setup connection followed by conversation**

Example TCP applications  
**Web, Email, Telnet**

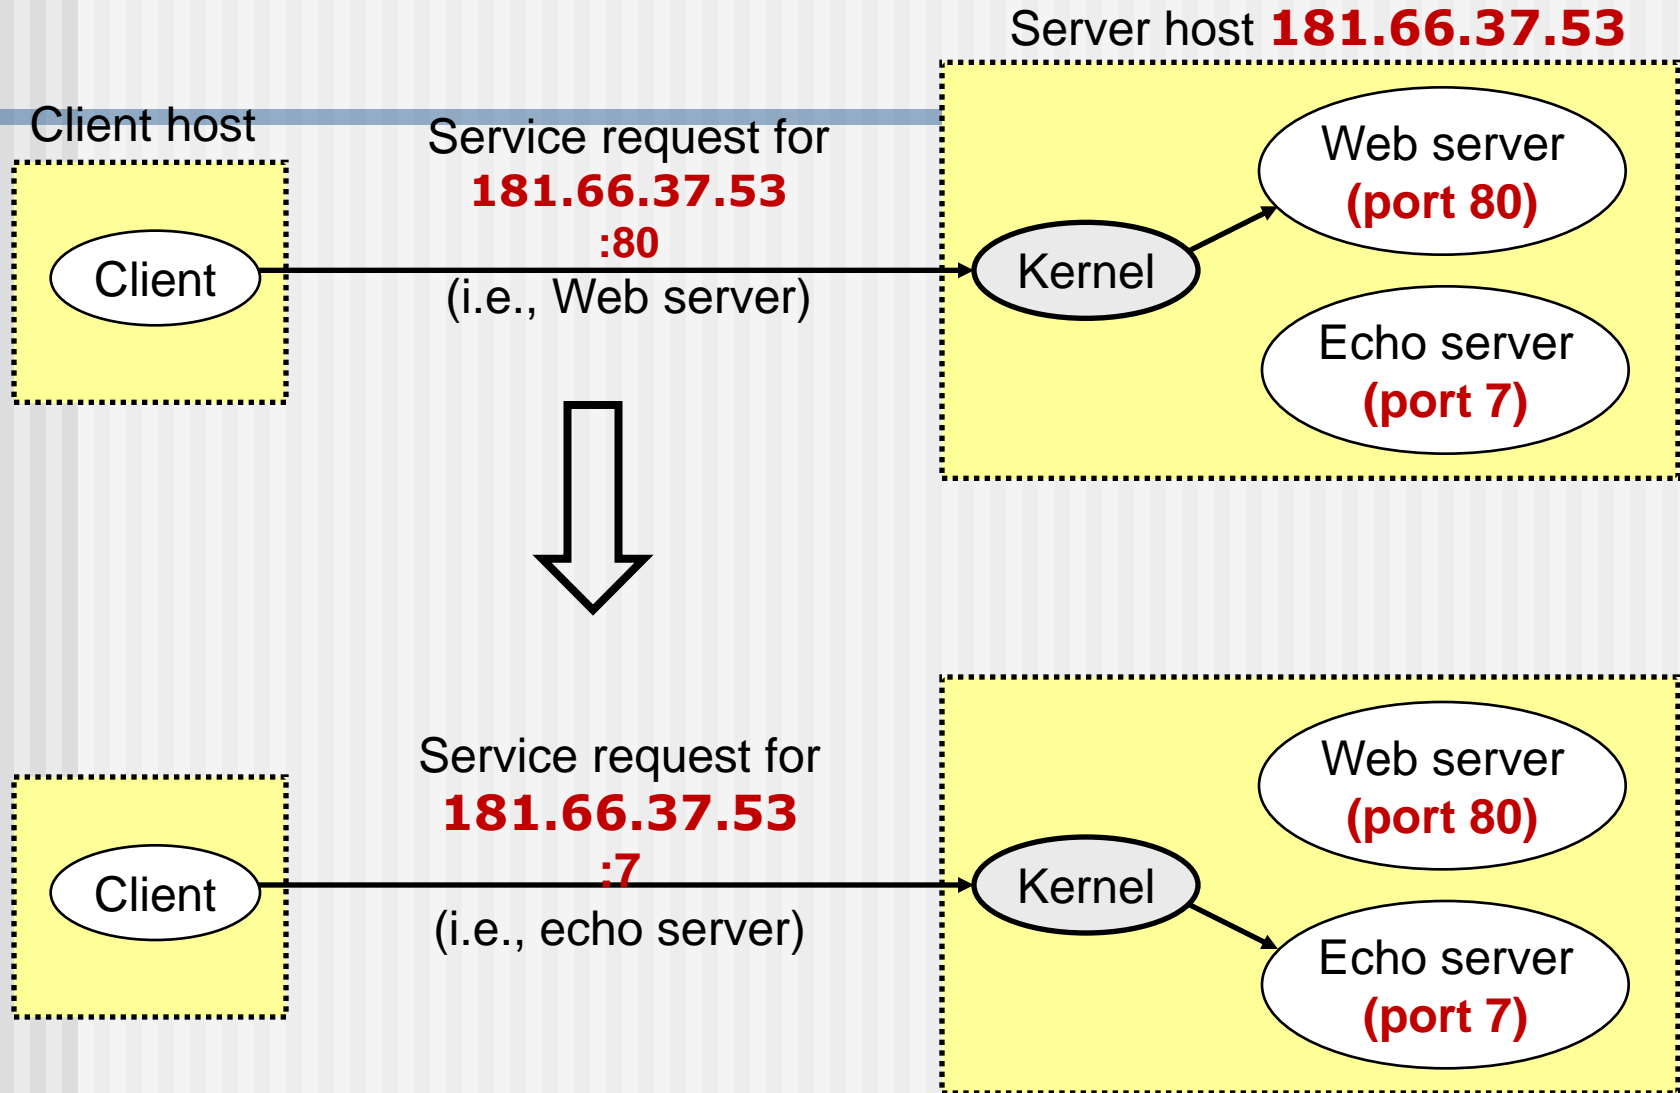
# PORTS (IPv4 vs. IPv6)

- **Port** - A 16-bit number that identifies the application process that receives an incoming message.
- **Local hosts for IPv4** address is 127.0.0.1
- **Local hosts for IPv6** address is IPv6 is 0:0:0:0:0:0:0:1 **shortened form is ::1**
- Both are stored in an *IP address struct* of appropriate type(library functions in C usually hide them)
  - in\_addr for IPv4
  - in6\_addr for IPv6
- Port numbers divided into three categories
  - Well Known Ports 0-1023
  - Registered Ports 1024-49151 by the IANA (Internet Assigned Numbers Authority), and represent *second tier* common ports (socks (1080), WINS (1512), kermi (1649), https (443))
  - Dynamic/Private Ports 49152-65535 *ephemeral* ports, available for temporary client usage

# Well Known Ports- Reserved ports

- Reserved ports or well-known ports (0 to 1023)
  - Standard ports for well-known applications.
  - See `/etc/services` file on any UNIX machine for listing of services on reserved ports.
    - 1 TCP Port Service Multiplexer
    - 7 Echo Server
    - 20 File Transfer Protocol (FTP) Data
    - 21 FTP Control
    - 23 Telnet
    - 25 Simple Mail Transfer (SMT)
    - 43 Who Is
    - 69 Trivial File Transfer Protocol (TFTP)
    - 80 HTTP

# Using Ports to Identify Services



# The Socket----Socket as a File

- a **socket is like a file**: you can read/write to/from the network just like you would a file
- For connection-oriented communication (e.g. TCP)
  - servers (passive open) do listen and accept operations
  - clients (active open) do connect operations
  - both sides can then do read and/or write (or send and recv)
  - then each side must close etc
- Connectionless (e.g. UDP): uses sendto and recvfrom

# Sockets And Socket Libraries

---

- **In Unix, socket procedures** (e.g. listen, connect, etc.) **are *system calls***
  - part of the operating system
  - when you call the function, control moves to the operating system, and you are using “system” CPU time

# Sockets And Socket Libraries

---

- **On some operating systems, socket procedures are *not* part of the OS e.g in Windows**
  - instead, they are implemented as a library, linked into the application object code (e.g. a **DLL under Windows**)
  - Typically, this DLL makes calls to similar procedures that are part of the native operating system.

# The Most Popular Socket Interface

---

- The **Berkeley Sockets API**
  - Originally developed as part of BSD Unix
  - BSD = Berkeley Software Distribution
    - API=Application Program Interface
  - Now the most popular **API for C/C++** programmers writing applications over TCP/IP
    - Also emulated in other languages: Perl, Tcl/Tk, Python etc.
    - Also emulated on other operating **systems**: **Windows**, etc.



# Sockets and Data types

## ■ Data types

---

int8\_t  
int16\_t  
int32\_t

signed 8-bit integer  
signed 16-bit integer  
signed 32-bit integer

uint8\_t  
uint16\_t  
uint32\_t

unsigned 8-bit integer  
unsigned 16-bit integer  
unsigned 32-bit integer

**u\_char**

Unsigned 8-bit character

**u\_short**

Unsigned 16-bit integer

**u\_long**

Unsigned 32-bit integer


## More on Data types

---

- Internet Address Structure

```
struct in_addr  
{  
    in_addr_t    s_addr;  
};
```

```
struct in_addr  
{  
    u_long s_addr ;  
} ;
```

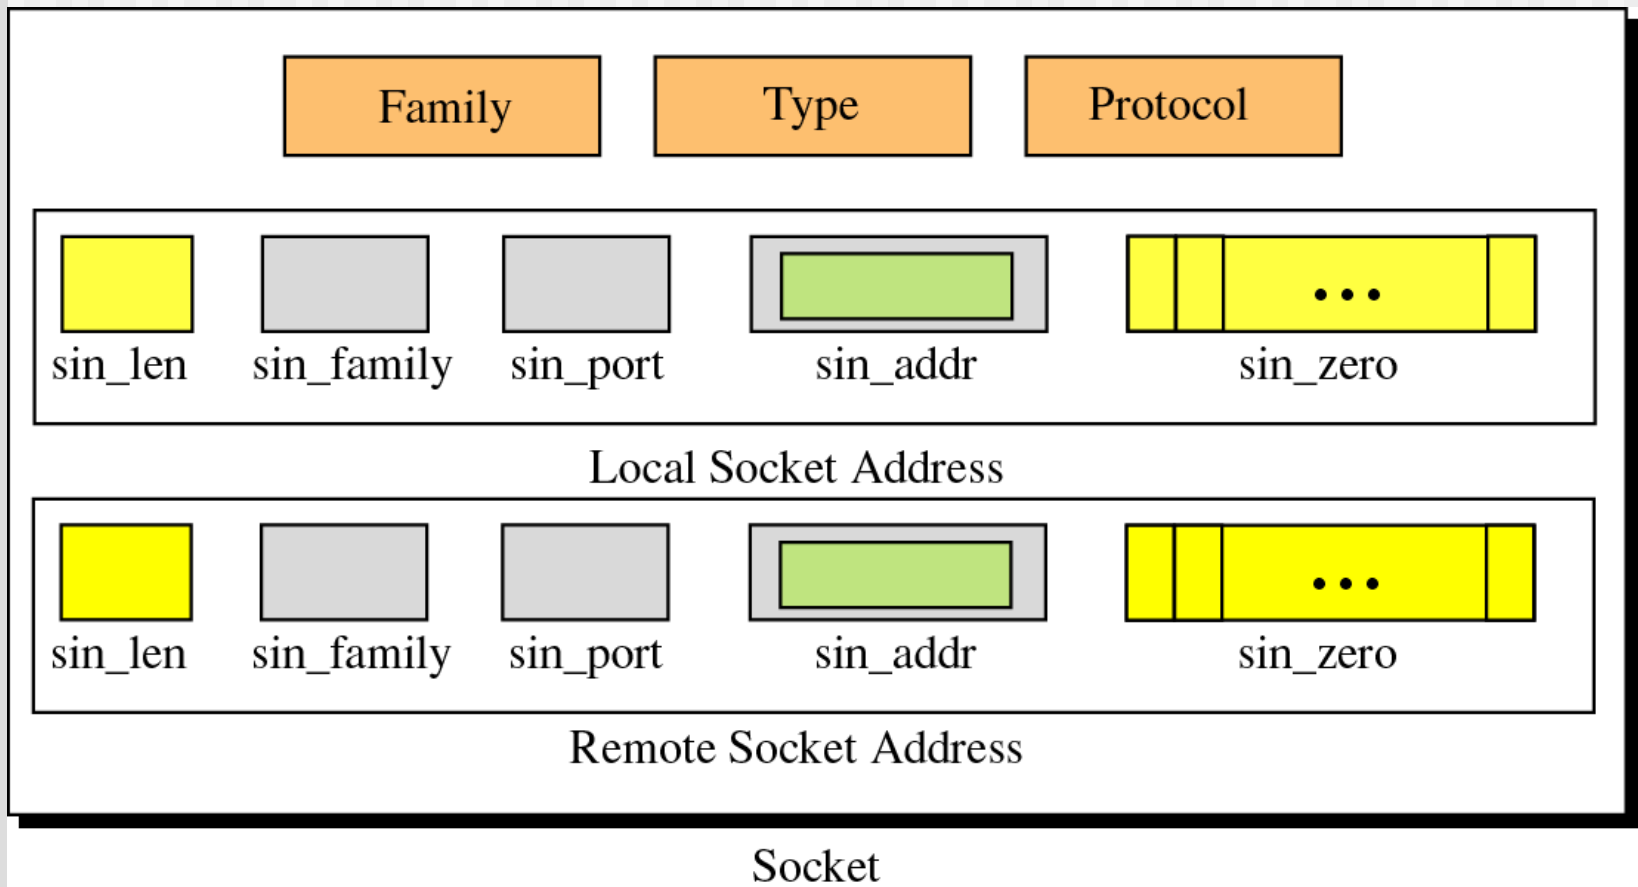


## sockaddr\_in

```

struct  sockaddr_in
{
    u_char           sin_len ;
    u_short          sin_family ;
    u_short          sin_port ;
    struct in_addr   sin_addr ;
    char             sin_zero [8] ;
};
  
```

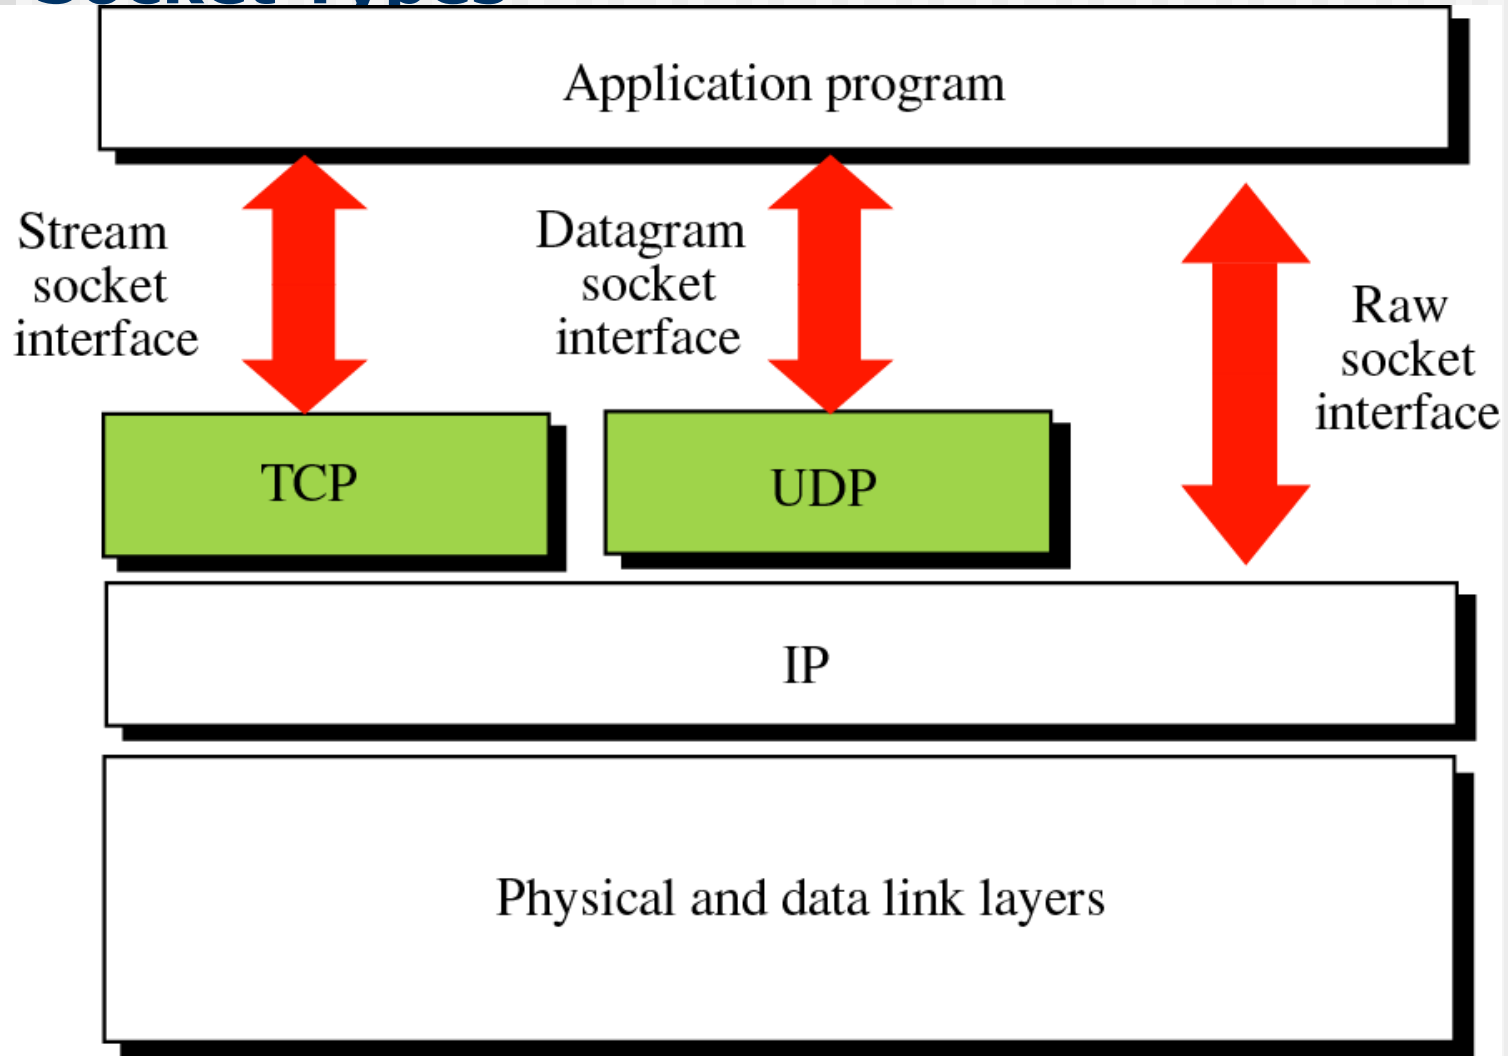
# Socket Structure



# Sockets

- Sockets provide a standard interface between network and application –they are Independent of network type:
- Commonly used with TCP/IP and UDP/IP,
- Two types of socket:
  - **Stream** – provides a virtual circuit service
  - **Datagram** – delivers individual packets

## Socket Types



# Byte Ordering


- **two** ways to store two bytes in memory: with the lower-order byte at the starting address (*little-endian* byte order) or with the high-order byte at the starting address (*big-endian* byte order). They are called *host byte order*.
- an Intel processor stores the **32-bit integer as four consecutives** bytes in memory in the order 1-2-3-4, where 1 is the **most significant byte**. IBM PowerPC processors would store the integer in the byte order 4-3-2-1.

## Little-Endian (Intel)



← **Address A** →  
← **Address A+1** →

## Big-Endian (RISC-Sparc)



# Byte Order and Networking

- Suppose a Big Endian machine sends a 16 bit integer with the value 2:

00000000000000010

- A Little Endian machine will understand the number as 512:

0000001000000000

- How do two machines with different byte-orders communicate?
  - Using network byte-order
  - Network byte-order = big-endian order



# network byte-order conversion

---

- The following functions are used for conversion to the network order
- **The htons(), htonl(), ntohs(), and ntohl() Functions**  
#include <netinet/in.h>

uint16\_t htons(uint16\_t host16bitvalue);

uint32\_t htonl(uint32\_t host32bitvalue);

uint16\_t ntohs(uint16\_t net16bitvalue);

uint32\_t ntohl(uint32\_t net32bitvalue);

- **first two** return the value in network byte order (16 and 32 bit, respectively). The **latter return** the value in host byte order (16 and 32 bit, respectively).

# Network Byte Order

---


- Conversion of application-level data is left up to the presentation layer.
- Lower level layers communicate using a fixed byte order called *network byte order* for all control data.
- TCP/IP mandates that *big-endian byte ordering* be used for transmitting protocol information
- All values stored in a `sockaddr_in` must be in network byte order.
  - `sin_port`      a TCP/IP port number.
  - `sin_addr`     an IP address.

# Network Byte Order Functions

---

- Several functions are provided to allow conversion between **host and network byte ordering**,
- Conversion macros (`<netinet/in.h>`)
  - **to translate 32-bit numbers (i.e. IP addresses):**
    - `unsigned long htonl(unsigned long hostlong);`
    - `unsigned long ntohl(unsigned long netlong);`
  - **to translate 16-bit numbers (i.e. Port numbers):**
    - `unsigned short htons(unsigned short hostshort);`
    - `unsigned short ntohs(unsigned short netshort);`

# Byte-Order Transformation



```
u_short  htons ( u_short  host_short );
```

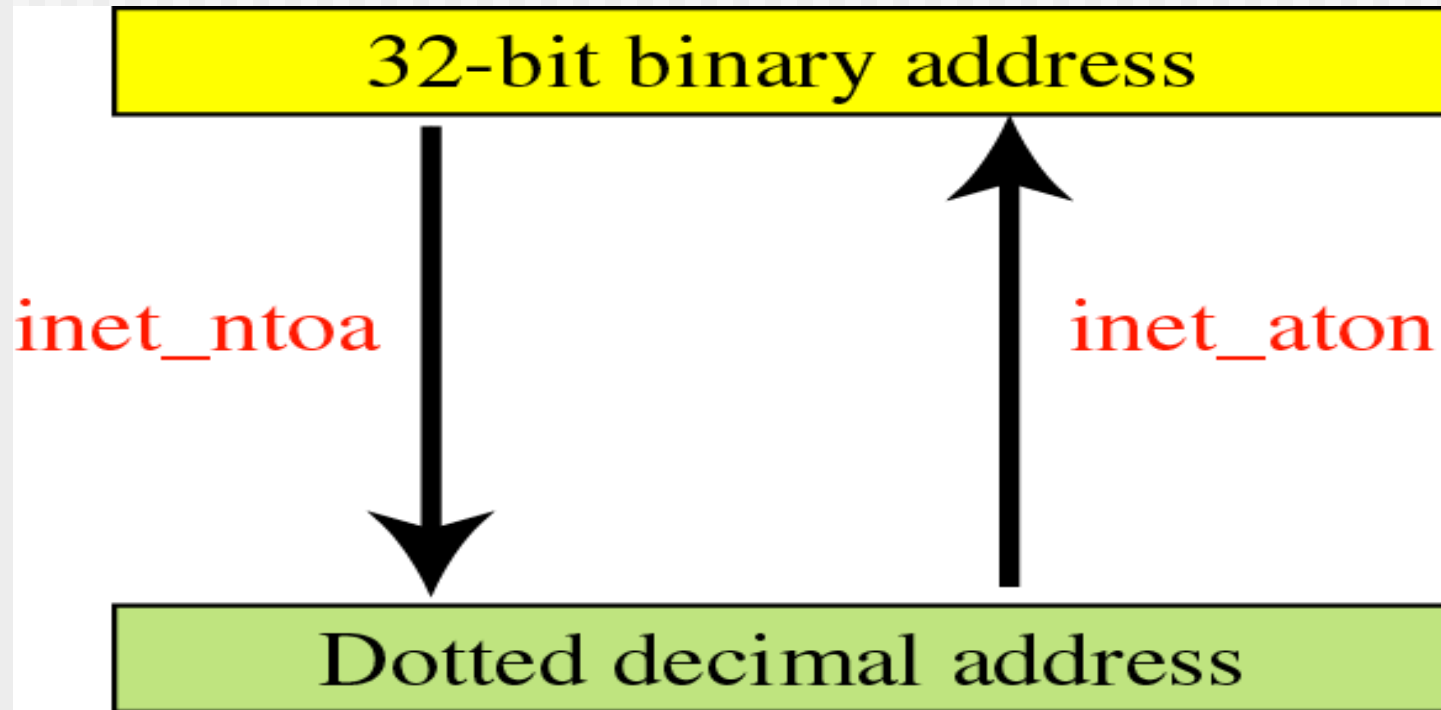
```
u_short  ntohs ( u_short  network_short );
```

```
u_long   htonl ( u_long   host_long );
```

```
u_long   ntohl ( u_long   network_long );
```

# Address Transformation

```
int      inet_aton ( const char  *strptr , struct in_addr *addrptr ) ;  
char     *inet_ntoa (struct in_addr inaddr ) ;
```



## Byte-Manipulation Functions

---

- In network programming, we often need to initialize a **field**, copy the contents of one field to another, or compare the contents of two fields.
  - Cannot use string functions (**strcpy**, **strcmp**, ...) which assume null character termination.

```
void *memset ( void *dest , int chr , int len ) ;
```

```
void *memcpy ( void *dest , const void *src , int len ) ;
```

```
int memcmp ( const void *first , const void *second , int len ) ;
```

# Creating a Socket: The `socket()` Function

```
#include <sys/types.h>
#include <sys/socket.h>
int socket(int domain, int type, int protocol);
```

- `domain` is one of the *Protocol Families* (`AF_INET`, `AF_UNIX`, etc.)
- `type` defines the communication protocol semantics, usually defines either:
  - `SOCK_STREAM`: connection-oriented stream (TCP)
  - `SOCK_DGRAM`: connectionless, unreliable (UDP)
- `protocol` specifies a particular protocol, just set this to 0 to accept the default
- Example `int sockfd = socket (AF_INET, SOCK_STREAM, 0);`

## Socket Primitives ---in the library

Primitive	Meaning
Socket	Create a new communication endpoint
Bind	Attach a local address to a socket
Listen	Announce willingness to accept connections
Accept	Block caller until a connection request arrives
Connect	Actively attempt to establish a connection
Send	Send some data over the connection
Receive	Receive some data over the connection
Close	Release the connection



# Creating a Socket-2:The socket() Function

- `int sockid = socket(family, type, protocol) ;`
  - **sockid**: socket descriptor, an integer (like a file-handle)
  - **family**: integer, communication domain, e.g.,
    - PF\_INET, IPv4 protocols, Internet addresses (typically used)
    - PF\_UNIX, Local communication, File addresses
  - **type**: communication type
    - SOCK\_STREAM - reliable, 2-way, connection-based service
    - SOCK\_DGRAM - unreliable, connectionless, messages of maximum length
  - **protocol**: specifies protocol
    - IPPROTO\_TCP IPPROTO\_UDP
    - usually set to 0 (i.e., use default protocol)
  - upon failure returns -1
- ☞ NOTE: socket call does not specify where data will be coming from, nor where it will be going to – it just creates the interface!

# Programming Sockets

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int fd;
...
fd = socket(family, type, protocol);
if (fd == -1) {
    // Error: unable to create socket
    ...
}
...
```

AF\_INET for IPv4  
AF\_INET6 for IPv6

SOCK\_STREAM for TCP  
SOCK\_DGRAM for UDP

0 (not used for Internet sockets)


# Specifying Addresses

- Socket API defines a **generic** data type for addresses:

```
struct sockaddr {  
    unsigned short sa_family; /* Address family (e.g. AF_INET) */  
    char sa_data[14];        /* Family-specific address information */  
}
```

- Particular form of the sockaddr used for **TCP/IP** addresses:

```
struct in_addr {  
    unsigned long s_addr; /* Internet address (32 bits) */  
}  
  
struct sockaddr_in {  
    unsigned short sin_family; /* Internet protocol (AF_INET) */  
    unsigned short sin_port; /* Address port (16 bits) */  
    struct in_addr sin_addr; /* Internet address (32 bits) */  
    char sin_zero[8]; /* Not used */  
}
```

 **Important:** sockaddr\_in can be casted to a sockaddr

# The connect() Function

---

- The **connect()** function is used by a TCP client to establish a connection with a TCP server

```
#include <sys/socket.h>
```

```
int connect (int sockfd, const struct sockaddr *servadr,  
            socklen_t addrlen);
```

- where sockfd is the socket descriptor returned by the socket function.
- The function **returns 0** if it succeeds in establishing a connection (i.e., successful TCP three-way handshake, **-1 otherwise**).

# The bind() Function

---

- The bind() assigns a local protocol address to a socket.

```
#include <sys/socket.h>
int bind(int sockfd, const struct sockaddr *servaddr, socklen_t addrlen);
```

where **sockfd** is the socket descriptor, **servaddr** is a pointer to a protocol-specific address and **addrlen** is the size of the address structure.

bind() returns 0 if it succeeds, -1 on error.

## The bind() Function-2 example in IPV4

- use of generic socket address `sockaddr` requires that any calls to these functions **must cast** the pointer to the protocol-specific address structure

```
struct sockaddr_in serv; /* IPv4 socket address structure */
```

```
bind(sockfd, (struct sockaddr*) &serv, sizeof(serv))
```

- **A process can bind a specific IP address to its socket: for a TCP client, this assigns the source IP address that will be used for IP datagrams sent on the sockets. For a TCP server, this restricts the socket to receive incoming client connections destined only to that IP address.**
- TCP client does not bind an IP address to its socket. The kernel chooses the **source IP socket** to be connected, based on the outgoing interface that is used. If a TCP server does not bind an IP address to its socket, the kernel uses the destination IP address of the incoming packets as the server's source address.

## The listen() Function

The listen() function converts an unconnected socket into a passive socket, indicating that the kernel should accept incoming connection requests directed to this socket.

```
#define _OE_SOCKETS
#include <sys/socket.h>
int listen(int socket, int backlog);
```

*Where*

*Backlog* Defines the maximum length for the queue of pending connections.

The function listen() return 0 if it succeeds, -1 on error.

# The accept() Function

---

- used by a server to accept a connection request from a client

```
#include <sys/socket.h>
```

```
int accept(int sockfd, struct sockaddr *cliaddr,  
socklen_t *addrlen);
```

**sockfd** is a new file descriptor that is connected to the client that called the connect()

**cliaddr and addrlen** arguments are used to return the protocol address of the client



# The send() Function

---

- Remember Socket as a file ? We can use *read* and *write* to communicate with a socket as long as it is connected.
- If we need to specify other options , other than read and write plainly , we use other functions. E.g here send() is similar to write() but allows to specify some options. *buf* and *nbytes* have the same meaning as they have with *write*. *flags* is used to specify how we want the data to be transmitted

```
#include <sys/socket.h>
ssize_t send(int sockfd, const void *buf, size_t nbytes, i
nt flags);
```

**The function returns the number of bytes if it succeeds, -1 on error**

# The receive() Function

---

- The `recv()` function is similar to `read()`, but allows to specify some options

```
#include <sys/socket.h>  
ssize_t recv(int sockfd, void *buf, size_t nbytes, int flags);
```

**The function returns the length of the message in bytes, 0 if no messages are available and peer had done an orderly shutdown, or -1 on error.**

## The close() Function

---

- close() function is used to close a socket and terminate a TCP socket. It returns 0 if it succeeds, -1 on error.

```
#include <unistd.h>
```

```
int close(int sockfd);
```

## More on Networking and Sockets

---

- functions ***getaddrinfo()*** and ***getnameinfo()*** convert domain names, hostnames, and IP addresses between human-readable text representations and structured binary formats for the operating system's networking API.

# Working Exercises

---

- To be shared in the Piazza

# Sources

---

- <https://beej.us/guide/bgnet/html//index.html>
- <https://academy.nordicsemi.com/courses/cellular-iot-fundamentals/lessons/lesson-3-cellular-fundamentals/topic/lesson-3-exercise-1/>
- <https://www.youtube.com/watch?v=WdE3PCHSBy8>
- <https://people.cs.rutgers.edu/~pxk/rutgers/notes/sockets/>